CS/ECE 3810: Computer Organization

Lecture 4: MIPS instruction set

Anton Burtsev
August, 2022

Instruction Set

- Understanding the language of the hardware is key to understanding the hardware/software interface

- A program (in say, C) is compiled into an executable that is composed of machine instructions – this executable must also run on future machines – for example, each Intel processor reads in the same x86 instructions, but each processor handles instructions differently

- Java programs are converted into portable bytecode that is converted into machine instructions during execution (just-in-time compilation)

- What are important design principles when defining the instruction set architecture (ISA)?

Instruction Set

- Important design principles when defining the instruction set architecture (ISA):

  - keep the hardware simple – the chip must only implement basic primitives and run fast
  - keep the instructions regular – simplifies the decoding/scheduling of instructions

  We will later discuss RISC vs CISC

A Basic MIPS Instruction

C code: a = b + c ;

Assembly code: (human-friendly machine instructions)
                add   a, b, c   # a is the sum of b and c

A Basic MIPS Instruction

C code: a = b + c ;

Assembly code: (human-friendly machine instructions)
            add   a, b, c   # a is the sum of b and c

Machine code: (hardware-friendly machine instructions)
        00000010001100100100000000100000

Translate the following C code into assembly code:
        a = b + c + d + e;

Example

C code   a = b + c + d + e;
translates into the following assembly code:

Example

C code   a = b + c + d + e;

translates into the following assembly code:

```
add  a, b, c            add  a, b, c
add  a, a, d      or    add  f, d, e
add  a, a, e            add  a, a, f
```

- Instructions are simple: fixed number of operands (unlike C)
- A single line of C code is converted into multiple lines of assembly code
- Some sequences are better than others… the second sequence needs one more (temporary) variable f

Subtract Example

C code   f = (g + h) − (i + j);

Assembly code translation with only add and sub instructions:

Subtract Example

C code   f = (g + h) − (i + j);
translates into the following assembly code:

```
add  t0, g, h           add  f, g, h
add  t1,  i, j     or    sub   f, f, i
sub  f,   t0, t1         sub   f, f, j
```

Operands

- In C, each "variable" is a location in memory

- In hardware, each memory access is expensive – if variable *a* is accessed repeatedly, it helps to bring the variable into an on-chip scratchpad and operate on the scratchpad (registers)

- To simplify the instructions, we require that each instruction (add, sub) only operate on registers

- Note: the number of operands (variables) in a C program is very large; the number of operands in assembly is fixed… there can be only so many scratchpad registers

Registers

- The MIPS ISA has 32 registers (x86 has 8 registers) – Why not more? Why not less?

- Each register is 32 bits wide (modern 64-bit architectures have 64-bit wide registers)

- A 32-bit entity (4 bytes) is referred to as a word

- To make the code more readable, registers are partitioned as $s0-$s7 (C/Java variables), $t0-$t9 (temporary variables)…

Binary Stuff

- 8 bits = 1 Byte, also written as 8b = 1B

- 1 word = 32 bits = 4B

- 1KB = 1024 B = 2^{10} B

- 1MB = 1024 x 1024 B = 2^{20} B

- 1GB = 1024 x 1024 x 1024 B = 2^{30} B

- A 32-bit memory address refers to a number between 0 and $2^{32} - 1$, i.e., it identifies a byte in a 4GB memory

Memory Operands

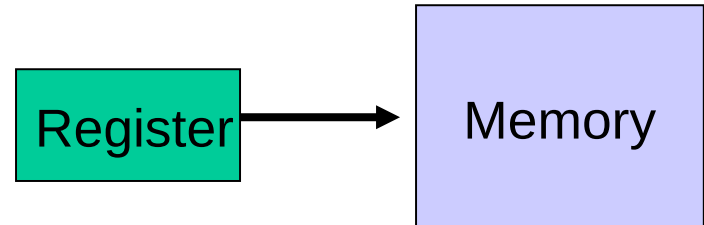- Values must be fetched from memory before (add and sub) instructions can operate on them

Load word
lw $t0, memory-address

```
Register ← Memory
```
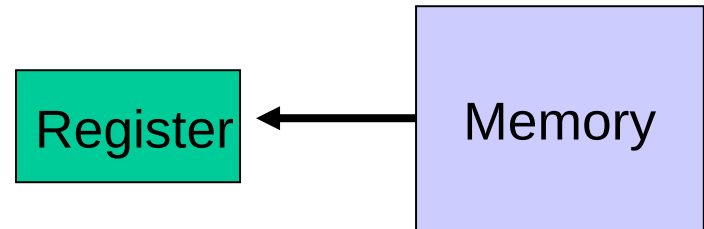
Store word
sw $t0, memory-address

```
Register → Memory
```
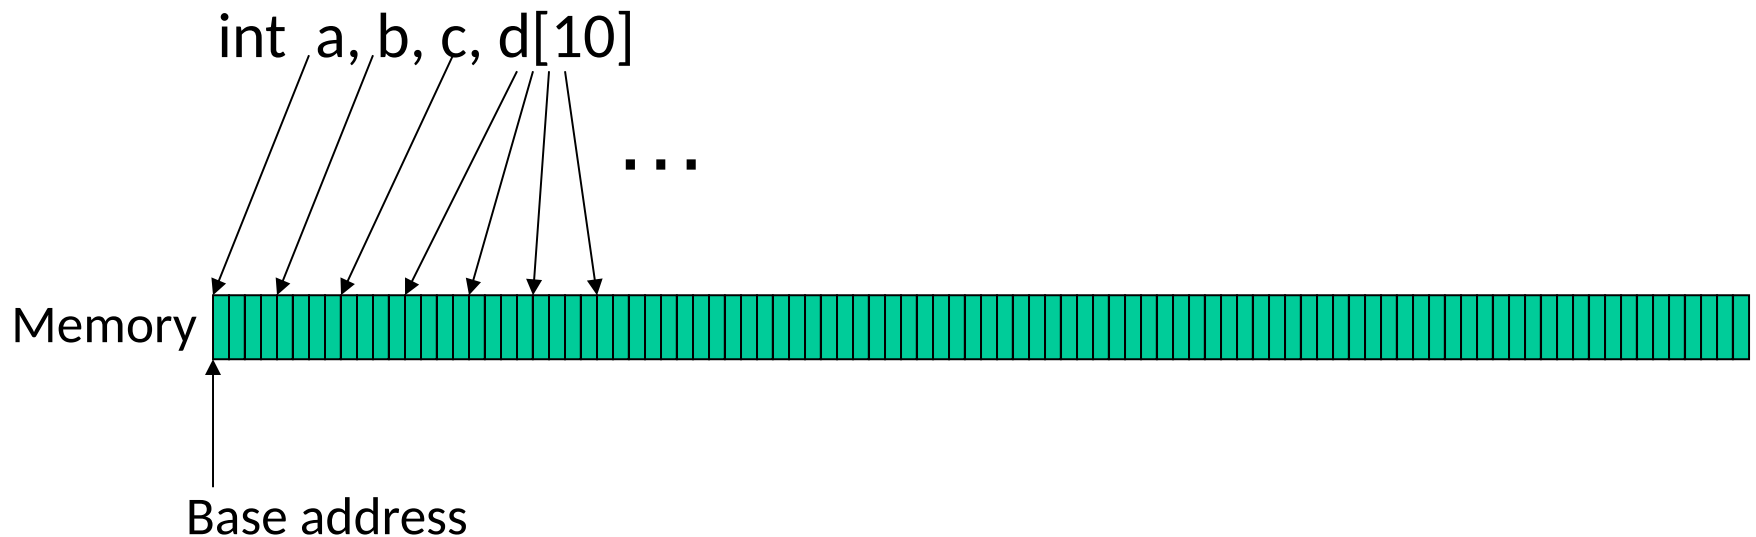
How is memory-address determined?

Memory Address

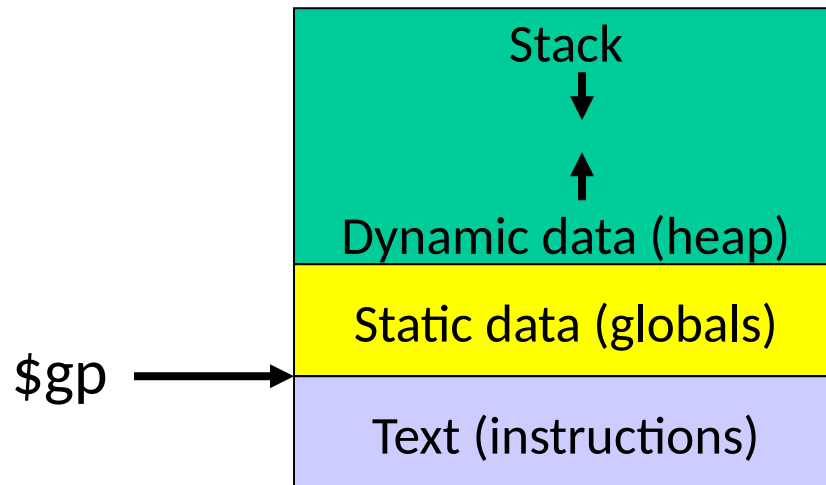- The compiler organizes data in memory… it knows the location of every variable (saved in a table)… it can fill in the appropriate mem-address for load-store instructions

int a, b, c, d[10]

…

Memory

Base address

Memory Organization

$gp points to area in memory that saves global variables

$gp →

Stack ↓ ↑
Dynamic data (heap)
Static data (globals)
Text (instructions)

Memory Instruction Format

- The format of a load instruction:

destination register

source address

lw $t0,   8($t3)

any register

a constant that is added to the register in parentheses
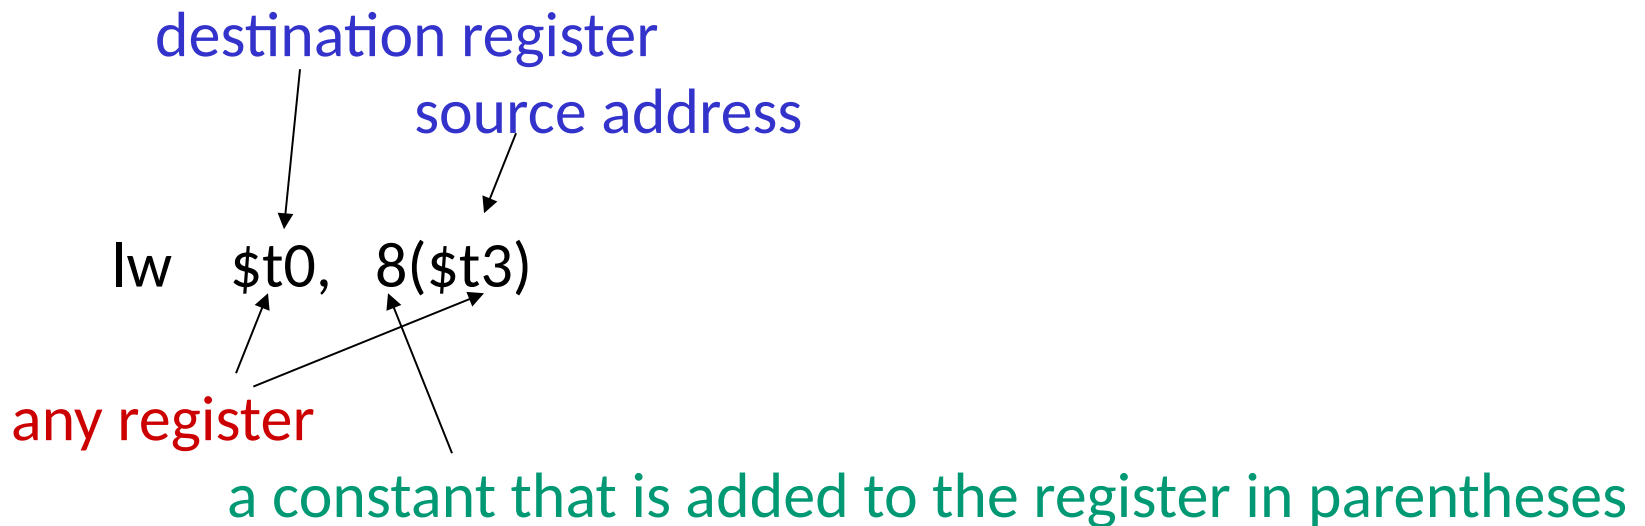
Memory Instruction Format

- The format of a store instruction:

source register

destination address

sw $t0,   8($t3)

any register

a constant that is added to the register in parentheses
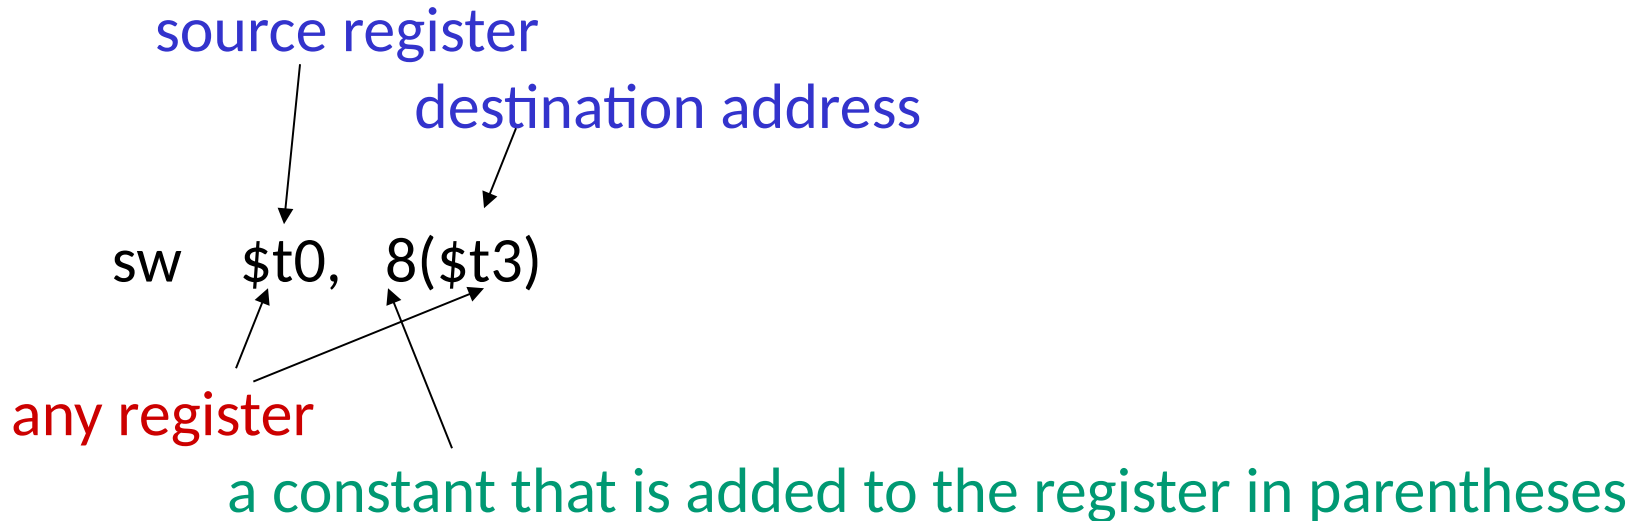
Example

int a, b, c, d[10];


```
addi   $gp, $zero, 1000   # assume that data is stored at
                          # base address 1000; placed in $gp;
                          # $zero is a register that always
                          # equals zero
lw   $s1, 0($gp)          # brings value of a into register $s1
lw   $s2, 4($gp)          # brings value of b into register $s2
lw   $s3, 8($gp)          # brings value of c into register $s3
lw   $s4, 12($gp)         # brings value of d[0] into register $s4
lw   $s5, 16($gp)         # brings value of d[1] into register $s5
```

Example

Convert to assembly:

Remember: int a, b, c, d[10];

C code: d[3] = d[2] + a;

Example

Convert to assembly:

Remember: int a, b, c, d[10];

C code: d[3] = d[2] + a;

Assembly (same assumptions as previous example):

```
lw    $s0, 0($gp)    #  a is brought into $s0
lw    $s1, 20($gp)   #  d[2] is brought into $s1
add   $s2, $s0, $s1  #  the sum is in $s2
sw    $s2, 24($gp)   #  $s2 is stored into d[3]
```

Assembly version of the code continues to expand!

Memory Organization

- The space allocated on stack by a procedure is termed the activation record (includes saved values and data local to the procedure) – frame pointer points to the start of the record and stack pointer points to the end – variable addresses are specified relative to $fp as $sp may change during the execution of the procedure
- $gp points to area in memory that saves global variables
- Dynamically allocated storage (with malloc()) is placed on the heap

Stack

↓

↑

Dynamic data (heap)

Static data (globals)

Text (instructions)
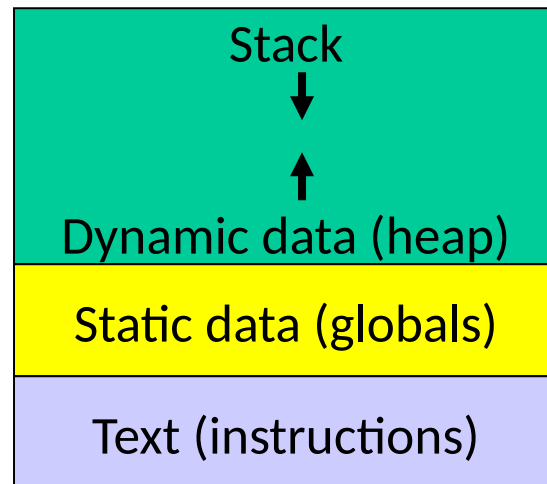
Binary Representation

- The binary number

01011000 00010101 00101110 11100111

Most significant bit   Least significant bit

represents the quantity
$0 \times 2^{31} + 1 \times 2^{30} + 0 \times 2^{29} + \ldots + 1 \times 2^{0}$

- A 32-bit word can represent 2^{32} numbers between
  0 and $2^{32}-1$
  … this is known as the unsigned representation as
    we're assuming that numbers are always positive

Negative Numbers

32 bits can only represent 2^{32} numbers – if we wish to also represent negative numbers, we can represent 2^{31} positive numbers (incl zero) and 2^{31} negative numbers

0000 0000 0000 0000 0000 0000 0000 0000$_{two}$ = 0$_{ten}$

0000 0000 0000 0000 0000 0000 0000 0001$_{two}$ = 1$_{ten}$

...

0111 1111 1111 1111 1111 1111 1111 1111$_{two}$ = 2^{31}-1

1000 0000 0000 0000 0000 0000 0000 0000$_{two}$ = -2^{31}

1000 0000 0000 0000 0000 0000 0000 0001$_{two}$ = -(2^{31} – 1)

1000 0000 0000 0000 0000 0000 0000 0010$_{two}$ = -(2^{31} – 2)

...

1111 1111 1111 1111 1111 1111 1111 1110$_{two}$ = -2

1111 1111 1111 1111 1111 1111 1111 1111$_{two}$ = -1

23

2's Complement

0000 0000 0000 0000 0000 0000 0000 0000$_\text{two}$ = 0$_\text{ten}$

0000 0000 0000 0000 0000 0000 0000 0001$_\text{two}$ = 1$_\text{ten}$

...

0111 1111 1111 1111 1111 1111 1111 1111$_\text{two}$ = 2^{31}-1

1000 0000 0000 0000 0000 0000 0000 0000$_\text{two}$ = -2^{31}

1000 0000 0000 0000 0000 0000 0000 0001$_\text{two}$ = -(2^{31} – 1)

1000 0000 0000 0000 0000 0000 0000 0010$_\text{two}$ = -(2^{31} – 2)

...

1111 1111 1111 1111 1111 1111 1111 1110$_\text{two}$ = -2

1111 1111 1111 1111 1111 1111 1111 1111$_\text{two}$ = -1

Why is this representation favorable?

Consider the sum of 1 and -2 .... we get -1

Consider the sum of 2 and -1 .... we get +1

This format can directly undergo addition without any conversions!

Each number represents the quantity

x_{31} -2^{31} + x_{30} 2^{30} + x_{29} 2^{29} + ... + x_1 2^1 + x_0 2^0

2's Complement

0000 0000 0000 0000 0000 0000 0000 0000$_{two}$ = 0$_{ten}$

0000 0000 0000 0000 0000 0000 0000 0001$_{two}$ = 1$_{ten}$

...

0111 1111 1111 1111 1111 1111 1111 1111$_{two}$ = 2^{31}-1

1000 0000 0000 0000 0000 0000 0000 0000$_{two}$ = -2^{31}

1000 0000 0000 0000 0000 0000 0000 0001$_{two}$ = -(2^{31} – 1)

1000 0000 0000 0000 0000 0000 0000 0010$_{two}$ = -(2^{31} – 2)

...

1111 1111 1111 1111 1111 1111 1111 1110$_{two}$ = -2

1111 1111 1111 1111 1111 1111 1111 1111$_{two}$ = -1

Note that the sum of a number x and its inverted representation x' always equals a string of 1s (-1).

$x + x' = -1$

$x' + 1 = -x$ … hence, can compute the negative of a number by

$-x = x' + 1$ inverting all bits and adding 1

Similarly, the sum of x and –x gives us all zeroes, with a carry of 1

In reality, $x + (-x) = 2^n$ … hence the name 2's complement

25

Example

- Compute the 32-bit 2's complement representations for the following decimal numbers:
    5, -5, -6

Example

- Compute the 32-bit 2's complement representations for the following decimal numbers:
    5, -5, -6

  5:   0000 0000 0000 0000 0000 0000 0000 0101
  -5:   1111 1111 1111 1111 1111 1111 1111 1011
  -6:   1111 1111 1111 1111 1111 1111 1111 1010

  Given -5, verify that negating and adding 1 yields the number 5

Recap – Numeric Representations

- Decimal $\qquad 35_{10} = 3 \times 10^1 + 5 \times 10^0$

- Binary $\qquad 00100011_2 = 1 \times 2^5 + 1 \times 2^1 + 1 \times 2^0$

- Hexadecimal (compact representation)

  $\qquad$ 0x 23 $\quad$ or $\quad 23_{hex} = 2 \times 16^1 + 3 \times 16^0$

0-15 (decimal) ✉ 0-9, a-f (hex)

Dec	Binary	Hex	Dec	Binary	Hex	Dec	Binary	Hex	Dec	Binary	Hex
0	0000	00	4	0100	04	8	1000	08	12	1100	0c
1	0001	01	5	0101	05	9	1001	09	13	1101	0d
2	0010	02	6	0110	06	10	1010	0a	14	1110	0e
3	0011	03	7	0111	07	11	1011	0b	15	1111	0f

Constant or Immediate Operands

- We often use constants in operations

- Example: add 4 to register $s3

  ```
  lw $t0, AddrConstant4($s1)# $t0 = constant 4
  add $s3,$s3,$t0 # $s3 = $s3 + $t0 ($t0 == 4)
  ```

- A more elegant way

  ```
  addi $s3,$s3,4 # $s3 = $s3 + 4
  ```

Instruction format (R-type)

- Instructions are 32bit words in memory

op	rs	rt	rd	shamt	funct
6 bits	5 bits	5 bits	5 bits	5 bits	6 bits

- *op*: Basic operation of the instruction, traditionally called the opcode
- *rs*: The first register source operand
- *rt*: The second register source operand
- *rd*: The register destination operand. It gets the result of the operation
- *shamt*: Shift amount
- *funct*: Function/function code, selects the specific variant of the operation in the op field

Instruction format (R-type)

- Instructions are 32bit words in memory

op	rs	rt	rd	shamt	funct
6 bits	5 bits	5 bits	5 bits	5 bits	6 bits

- This works ok well for instructions like
  - add $s0, $s1, $s3
- But what about
  - lw $t0, 32($s0)
  - addi $t0, $t1, 4   # t0 = t1 + 4

Instruction format (I-type)

- Instructions are 32bit words in memory
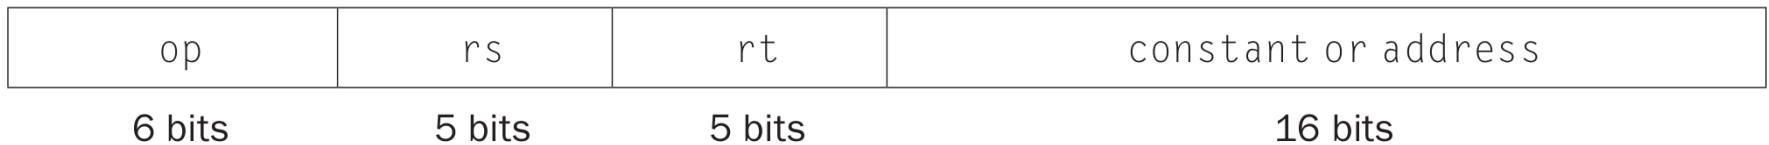
op	rs	rt	constant or address
6 bits	5 bits	5 bits	16 bits

- *op*: Basic operation of the instruction, traditionally called the opcode
- *rs*: The first register source operand
- *rt*: New meaning – destination register

Examples

Instruction	Format	op	rs	rt	rd	shamt	funct	address
add	R	0	reg	reg	reg	0	32_{ten}	n.a.
sub (subtract)	R	0	reg	reg	reg	0	34_{ten}	n.a.
add immediate	I	8_{ten}	reg	reg	n.a.	n.a.	n.a.	constant
lw (load word)	I	35_{ten}	reg	reg	n.a.	n.a.	n.a.	address
sw (store word)	I	43_{ten}	reg	reg	n.a.	n.a.	n.a.	address

FIGURE 2.5   MIPS instruction encoding. In the table above, "reg" means a register number between 0 and 31, "address" means a 16-bit address, and "n.a." (not applicable) means this field does not appear in this format. Note that add and sub instructions have the same value in the op field; the hardware uses the funct field to decide the variant of the operation: add (32) or subtract (34).

Register numbers

- $s0 - $s7 map on hardware registers 16 – 23
  - E.g., $s0 is 16, $s1 is 17
- $t0 - $t7 map on hardware registers 8 – 15
  - E.g., $t0 is 8, $t1 is 17

Example

```
A[300] = h + A[300]
```

- Gets compiled to

Example

A[300] = h + A[300]

- Gets compiled to

```
lw  $t0,1200($t1) # Temporary reg $t0 gets A[300]

add $t0,$s2,$t0   # Temporary reg $t0 gets h + A[300]

sw  $t0,1200($t1) # Stores h+A[300] back into A[300]
```

Example

A[300] = h + A[300]

- Gets compiled to

```
lw   $t0,1200($t1) # Temporary reg $t0 gets A[300]
add $t0,$s2,$t0    # Temporary reg $t0 gets h + A[300]
sw   $t0,1200($t1) # Stores h+A[300] back into A[300]
```

Op	rs	rt	rd	address/ shamt	funct
35	9	8		1200	
0	18	8	8	0	32
43	9	8		1200	

Instruction encoding

Instruction	Format	op	rs	rt	rd	shamt	funct	address
add	R	0	reg	reg	reg	0	32_{ten}	n.a.
sub (subtract)	R	0	reg	reg	reg	0	34_{ten}	n.a.
add immediate	I	8_{ten}	reg	reg	n.a.	n.a.	n.a.	constant
lw (load word)	I	35_{ten}	reg	reg	n.a.	n.a.	n.a.	address
sw (store word)	I	43_{ten}	reg	reg	n.a.	n.a.	n.a.	address

FIGURE 2.5   MIPS instruction encoding. In the table above, "reg" means a register number between 0 and 31, "address" means a 16-bit address, and "n.a." (not applicable) means this field does not appear in this format. Note that add and sub instructions have the same value in the op field; the hardware uses the funct field to decide the variant of the operation: add (32) or subtract (34).

Example

- Decimal

Op	rs	rt	rd	address/ shamt	funct
35	9	8	1200		
0	18	8	8	0	32
43	9	8	1200		

- Binary

100011	01001	01000	0000 0100 1011 0000		
000000	10010	01000	01000	00000	100000
101011	01001	01000	0000 0100 1011 0000		

MIPS machine language

Name	Format	Example						Comments
add	R	0	18	19	17	0	32	add $s1,$s2,$s3
sub	R	0	18	19	17	0	34	sub $s1,$s2,$s3
addi	I	8	18	17	100			addi $s1,$s2,100
lw	I	35	18	17	100			lw $s1,100($s2)
sw	I	43	18	17	100			sw $s1,100($s2)
Field size		6 bits	5 bits	5 bits	5 bits	5 bits	6 bits	All MIPS instructions are 32 bits long
R-format	R	op	rs	rt	rd	shamt	funct	Arithmetic instruction format
I-format	I	op	rs	rt	address			Data transfer format

FIGURE 2.6 MIPS architecture revealed through Section 2.5. The two MIPS instruction formats so far are R and I. The first 16 bits are the same: both contain an *op* field, giving the base operation; an *rs* field, giving one of the sources; and the *rt* field, which specifies the other source operand, except for load word, where it specifies the destination register. R-format divides the last 16 bits into an *rd* field, specifying the destination register; the *shamt* field, which Section 2.6 explains; and the *funct* field, which specifies the specific operation of R-format instructions. I-format combines the last 16 bits into a single *address* field.

Logical operations

Logical operations	C operators	Java operators	MIPS instructions
Shift left	<<	<<	sll
Shift right	>>	>>>	srl
Bit-by-bit AND	&	&	and, andi
Bit-by-bit OR	|	|	or, ori
Bit-by-bit NOT	~	~	nor

FIGURE 2.8   C and Java logical operators and their corresponding MIPS instructions. MIPS implements NOT using a NOR with one operand being zero.

Shift

- Shift left by 4

- Before
  - 0000 0000 0000 0000 0000 0000 0000 1001 = 9

- After
  - 0000 0000 0000 0000 0000 0000 1001 0000 = 144

Encoding shift

- Example

      sll $t2,$s0,4 # reg $t2 = reg $s0 << 4 bits

- Shift amount

op	rs	rt	rd	shamt	funct
0	0	16	10	4	0

Control instructions

Branch instructions

- Branch when equal

    `beq register1, register2, L1`

    – Go to `L1` if `register1` equals `register2`

- Branch when not equal

    `bne register1, register2, L1`

    – Go to `L1` if `register1` does *not* equal `register2`

- Unconditional jump

    – `j L1`

    – Jump to `L1`

If then .. else …

```
if (i == j) f = g + h; else f = g – h;
```

- Assume that f,g, h, i, and j are in $s0, $s1, etc.

If then .. else ...

```
if (i == j) f = g + h; else f = g – h;
```

- Assume that f,g, h, i, and j are in $s0, $s1, etc.

```
        bne $s3,$s4,Else  # go to Else if i ≠ j

        add $s0,$s1,$s2   # f = g + h (skipped if i ≠ j)

        j Exit            # unconditional jump to Exit

Else:

        sub $s0,$s1,$s2   # f = g – h (skipped if i = j)

Exit:
```

Loops

```
while (save[i] == k)

    i += 1;
```

- Assume that i and k are in $s3 and $s5

Loops

```
while (save[i] == k)

    i += 1;
```

- Assume that i and k are in $s3 and $s5

```
Loop: sll $t1,$s3,2

    add $t1,$t1,$s6    # $t1 = address of save[i]

    lw $t0,0($t1)      # Temp reg $t0 = save[i]

    bne $t0,$s5, Exit  # go to Exit if save[i] ≠ k

    addi $s3,$s3,1     # i = i + 1

    j Loop             # go to Loop

Exit:
```

Comparisons

- Set on less than

  ```
  slt $t0, $s3, $s4 # $t0 = 1 if $s3 < $s4
  ```

- Or with a constrant

  ```
  slti $t0,$s2,10 # $t0 = 1 if $s2 < 10
  ```

- Now you can use slt, slti, beq, and bne along with $zero (register that is always 0)

Branch instructions

- Branch when equal

  `beq register1, register2, L1`

  - Go to L1 if `register1` equals `register2`

- Branch when not equal

  `bne register1, register2, L1`

  - Go to L1 if `register1` does *not* equal `register2`

- Unconditional jump

  - `j L1`

  - Jump to L1

Branch instructions

- Branch when equal

  `beq register1, register2, L1`

  – Go to L1 if `register1` equals `register2`

- Branch when not equal

  `bne register1, register2, L1`

  – Go to L1 if `register1` does *not* equal `register2`

- Unconditional jump

  – `j L1`

  – Jump to L1

Signed / Unsigned

- The hardware recognizes two formats:

  unsigned (corresponding to the C declaration unsigned int)
  -- all numbers are positive, a 1 in the most significant bit
     just means it is a really large number

  signed (C declaration is signed int or just int)
  -- numbers can be +/- , a 1 in the MSB means the number
     is negative

This distinction enables us to represent twice as many
numbers when we're sure that we don't need negatives

Shift Operations

op	rs	rt	rd	shamt	funct
6 bits	5 bits	5 bits	5 bits	5 bits	6 bits

- shamt: how many positions to shift
- Shift left logical
  - Shift left and fill with 0 bits
  - `sll` by i bits multiplies by 2^i
- Shift right logical
  - Shift right and fill with 0 bits
  - `srl` by i bits divides by 2^i (unsigned only)

AND Operations

- Useful to mask bits in a word
  - Select some bits, clear others to 0

 `and $t0, $t1, $t2`

$t2	0000 0000 0000 0000 0000 1101 1100 0000
$t1	0000 0000 0000 0000 0011 1100 0000 0000
$t0	0000 0000 0000 0000 0000 1100 0000 0000

OR Operations

- Useful to include bits in a word
  - Set some bits to 1, leave others unchanged

`or $t0, $t1, $t2`

$t2	0000 0000 0000 0000 0000 1101 1100 0000

$t1	0000 0000 0000 0000 0011 1100 0000 0000

$t0	0000 0000 0000 0000 0011 1101 1100 0000

NOT Operations

- Useful to invert bits in a word
  - Change 0 to 1, and 1 to 0
- MIPS has NOR 3-operand instruction
  - a NOR b == NOT (a OR b)

`nor $t0, $t1, $zero` ← Register 0: always read as zero

$t1 | 0000 0000 0000 0000 0011 1100 0000 0000

$t0 | 1111 1111 1111 1111 1100 0011 1111 1111

Signed and unsigned comparisons

Consider a comparison instruction:
    slt $t0, $t1, $zero
and $t1 contains the 32-bit number   1111 01...01

What gets stored in $t0?

Signed and unsigned comparisons

Consider a comparison instruction:

    slt $t0, $t1, $zero

and $t1 contains the 32-bit number   1111 01...01

What gets stored in $t0?

The result depends on whether $t1 is a signed or unsigned
number – the compiler/programmer must track this and
accordingly use either slt or sltu

  slt $t0, $t1, $zero   stores 1 in $t0
  sltu $t0, $t1, $zero   stores 0 in $t0

Sign Extension

- Occasionally, 16-bit signed numbers must be converted into 32-bit signed numbers – for example, when doing an add with an immediate operand

- The conversion is simple: take the most significant bit and use it to fill up the additional bits on the left – known as sign extension

So 2_{10} goes from 0000 0000 0000 0010   to
0000 0000 0000 0000 0000 0000 0000 0010

and -2_{10} goes from 1111 1111 1111 1110   to
1111 1111 1111 1111 1111 1111 1111 1110

Procedures

Procedure Calling

■ Steps required

1. Place parameters in registers
2. Transfer control to procedure
3. Acquire storage for procedure
4. Perform procedure's operations
5. Place result in register for caller
6. Return to place of call

Calling functions

```
// some code...
foo();
// more code..
```

- $ra contains information for how to return from a subroutine
  - i.e., from foo()

- Functions can be called from different places in the program

```
if (a == 0) {
    foo();

    …

} else {

    foo();

    …

}
```

Procedure Call Instructions

- Procedure call: jump and link

  `jal ProcedureLabel`

  - Address of following instruction put in $ra
  - Jumps to target address

- Procedure return: jump register

  `jr $ra`

  - Copies $ra to program counter
  - Can also be used for computed jumps
    - e.g., for case/switch statements

Calling conventions

- Goal: re-entrant programs
    - How to pass arguments
        - On the stack?
        - In registers?
    - How to return values
        - On the stack?
        - In registers?
    - What registers have to be preserved
        - All? Some subset?
- Conventions differ from compiler, optimizations, etc.

Passing arguments

- First 4 arguments in registers
  - $a0 - $a3
- Other arguments on the stack
- Return values in registers
  - $v0 - $v1

Preserving registers

- $t0 – $t9: temporaries
  - Can be overwritten by callee

- $s0 – $s7: saved

  - Must be saved/restored by callee

Leaf Procedure Example

- C code:

```
int leaf_example (int g, h, i, j)
{ int f;
  f = (g + h) - (i + j);
  return f;
}
```

  - Arguments g, …, j in $a0, …, $a3
  - f in $s0 (hence, need to save $s0 on stack)
  - Result in $v0

Leaf Procedure Example

- MIPS code:

```
leaf_example:
    addi $sp, $sp, -4
    sw   $s0, 0($sp)
    add  $t0, $a0, $a1
    add  $t1, $a2, $a3
    sub  $s0, $t0, $t1
    add  $v0, $s0, $zero
    lw   $s0, 0($sp)
    addi $sp, $sp, 4
    jr   $ra
```

Save $s0 on stack

Procedure body

Result

Restore $s0

Return

Recursive invocations

```
foo(int a) {
    if (a == 0)
        return;
    a--;
    foo(a);
    return;
}

foo(4);
```

Non-Leaf Procedures

- Procedures that call other procedures

- For nested call, caller needs to save on the stack:

  - Its return address

  - Any arguments and temporaries needed after the call

- Restore from the stack after the call

Non-Leaf Procedure Example

- C code:

```
int fact (int n)
{
  if (n < 1) return f;
  else return n * fact(n - 1);
}
```

- Argument n in $a0
- Result in $v0

Non-Leaf Procedure Example

- MIPS code:

```
fact:
    addi $sp, $sp, -8      # adjust stack for 2 items
    sw   $ra, 4($sp)       # save return address
    sw   $a0, 0($sp)       # save argument
    slti $t0, $a0, 1       # test for n < 1
    beq  $t0, $zero, L1
    addi $v0, $zero, 1     # if so, result is 1
    addi $sp, $sp, 8       #   pop 2 items from stack
    jr   $ra               #   and return
L1: addi $a0, $a0, -1      # else decrement n
    jal  fact              # recursive call
    lw   $a0, 0($sp)       # restore original n
    lw   $ra, 4($sp)       #   and return address
    addi $sp, $sp, 8       # pop 2 items from stack
    mul  $v0, $a0, $v0     # multiply to get result
    jr   $ra               # and return
```

Local variables

What types of variables do you know?

- Or where these variables are allocated in memory?

What types of variables do you know?

- Global variables
  - Initialized → data section
  - Uninitalized → BSS
- Dynamic variables
  - Heap
- Local variables
  - Stack

Global variables

```
1. #include <stdio.h>
2.
3. char hello[] = "Hello";
4. int main(int ac, char **av)
5. {
6.     static char world[] = "world!";
7.     printf("%s %s\n", hello, world);
8.     return 0;
9. }
```

Global variables

```
1. #include <stdio.h>
2.
3. char hello[] = "Hello";
4. int main(int ac, char **av)
5. {
6.     static char world[] = "world!";
7.     printf("%s %s\n", hello, world);
8.     return 0;
9. }
```

- Allocated in the data section
  - It is split in initialized (non-zero), and non-initialized (zero)
  - As well as read/write, and read only data section

Global variables

Dynamic variables (heap)

```c
1. #include <stdio.h>
2. #include <string.h>
3. #include <stdlib.h>
4.
5. char hello[] = "Hello";
6. int main(int ac, char **av)
7. {
8.     char world[] = "world!";
9.     char *str = malloc(64);
10.    memcpy(str, "beautiful", 64);
11.    printf("%s %s %s\n", hello, str, world);
12.    return 0;
13. }
```

Dynamic variables (heap)

```
1. #include <stdio.h>
2. #include <string.h>
3. #include <stdlib.h>
4.
5. char hello[] = "Hello";
6. int main(int ac, char **av)
7. {
8.     char world[] = "world!";
9.     char *str = malloc(64);
10.    memcpy(str, "beautiful", 64);
11.    printf("%s %s %s\n", hello, str, world);
12.    return 0;
13.}
```

- Allocated on the heap
  - Special area of memory provided by the OS from where malloc() can allocate memory

Dynamic variables (heap)

Local variables

- Local variables

```c
1. #include <stdio.h>
2.
3. char hello[] = "Hello";
4. int main(int ac, char **av)
5. {
6.     //static char world[] = "world!";
7.     char world[] = "world!";
8.     printf("%s %s\n", hello, world);
9.     return 0;
10. }
```

Local variables...

- Each function has private instances of local variables

```
foo(int x) {
    int a, b, c;
    ...
    return;
}
```

- Function can be called recursively

```
foo(int x) {
    int a, b, c;
    a = x + 1;
    if ( a < 100 )
        foo(a);
    return;
}
```
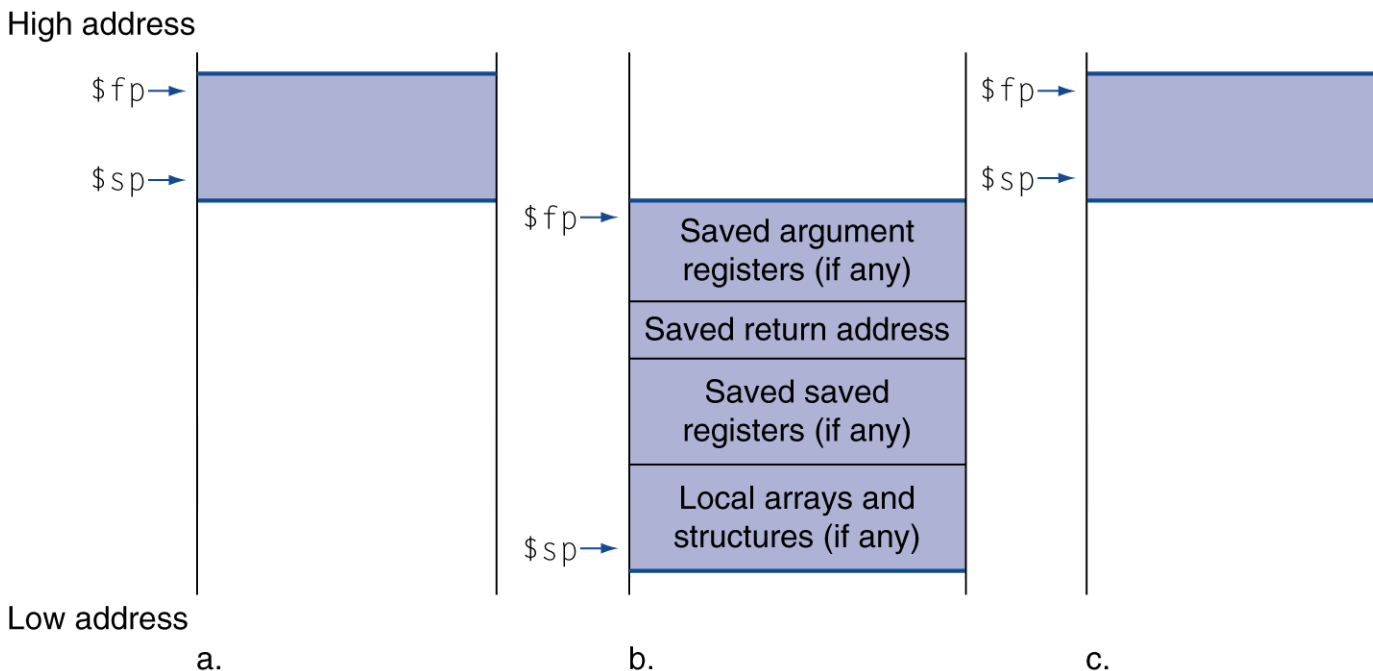
How to allocate local variables?

```
void my_function()
{
    int a, b, c;
    …
}
```

How to allocate local variables?

```
void my_function()
{
    int a, b, c;
    …
}
```
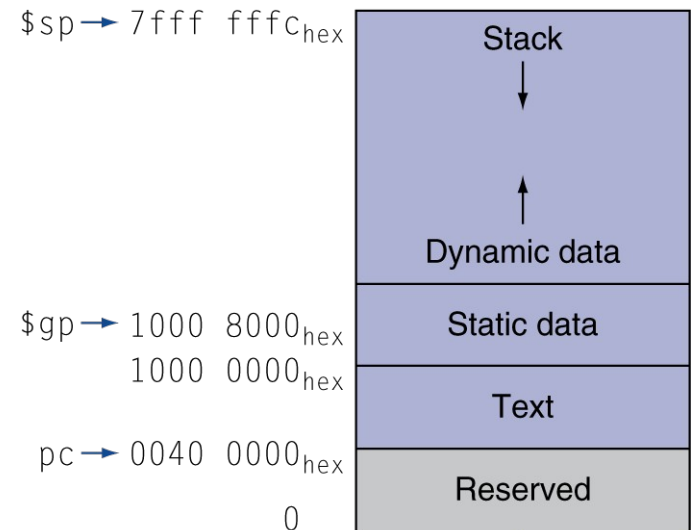
- On the stack!

Local Data on the Stack

High address

$fp →
$sp →

$fp →

Saved argument registers (if any)

Saved return address

Saved saved registers (if any)

Local arrays and structures (if any)

$sp →

$fp →
$sp →

Low address

a.   b.   c.

- Local data allocated by callee
  - e.g., C automatic variables
- Procedure frame (activation record)
  - Used by some compilers to manage stack storage

Memory Layout

- Text: program code
- Static data: global variables
  - e.g., static variables in C, constant arrays and strings
  - $gp initialized to address allowing ±offsets into this segment
- Dynamic data: heap
  - E.g., malloc in C, new in Java
- Stack: automatic storage

$sp → 7fff fffc_{hex}

$gp → 1000 8000_{hex}
1000 0000_{hex}

pc → 0040 0000_{hex}

0

| Stack |
| Dynamic data |
| Static data |
| Text |
| Reserved |

Character Data

- Byte-encoded character sets
  - ASCII: 128 characters
    - 95 graphic, 33 control
  - Latin-1: 256 characters
    - ASCII, +96 more graphic characters

- Unicode: 32-bit character set
  - Used in Java, C++ wide characters, …
  - Most of the world's alphabets, plus symbols
  - UTF-8, UTF-16: variable-length encodings

Byte/Halfword Operations

- Could use bitwise operations
- MIPS byte/halfword load/store
  - String processing is a common case

```
lb rt, offset(rs)        lh rt, offset(rs)
```

  - Sign extend to 32 bits in rt

```
lbu rt, offset(rs)       lhu rt, offset(rs)
```

  - Zero extend to 32 bits in rt

```
sb rt, offset(rs)        sh rt, offset(rs)
```

  - Store just rightmost byte/halfword

String Copy Example

- C code (naïve):
  - Null-terminated string

```
void strcpy (char x[], char y[])
{ int i;
  i = 0;
  while ((x[i]=y[i])!='\0')
    i += 1;
}
```

  - Addresses of x, y in $a0, $a1
  - i in $s0

String Copy Example

- MIPS code:

```
strcpy:
    addi $sp, $sp, -4       # adjust stack for 1 item
    sw   $s0, 0($sp)        # save $s0
    add  $s0, $zero, $zero  # i = 0
L1: add  $t1, $s0, $a1      # addr of y[i] in $t1
    lbu  $t2, 0($t1)        # $t2 = y[i]
    add  $t3, $s0, $a0      # addr of x[i] in $t3
    sb   $t2, 0($t3)        # x[i] = y[i]
    beq  $t2, $zero, L2     # exit loop if y[i] == 0
    addi $s0, $s0, 1        # i = i + 1
    j    L1                 # next iteration of loop
L2: lw   $s0, 0($sp)        # restore saved $s0
    addi $sp, $sp, 4        # pop 1 item from stack
    jr   $ra                # and return
```

32-bit Constants

- Most constants are small
  - 16-bit immediate is sufficient
- For the occasional 32-bit constant

```
lui rt, constant
```

  - Copies 16-bit constant to left 16 bits of rt
  - Clears right 16 bits of rt to 0

`lhi $s0, 61`

0000 0000 0111 1101	0000 0000 0000 0000

`ori $s0, $s0, 2304`

0000 0000 0111 1101	0000 1001 0000 0000

Branch Addressing

- Branch instructions specify
  - Opcode, two registers, target address
- Most branch targets are near branch
  - Forward or backward

op	rs	rt	constant or address
6 bits	5 bits	5 bits	16 bits

- PC-relative addressing
  - Target address = PC + offset × 4
  - PC already incremented by 4 by this time

Jump Addressing

- Jump (`j and jal`) targets could be anywhere in text segment
    - Encode full address in instruction

op	address
6 bits	26 bits
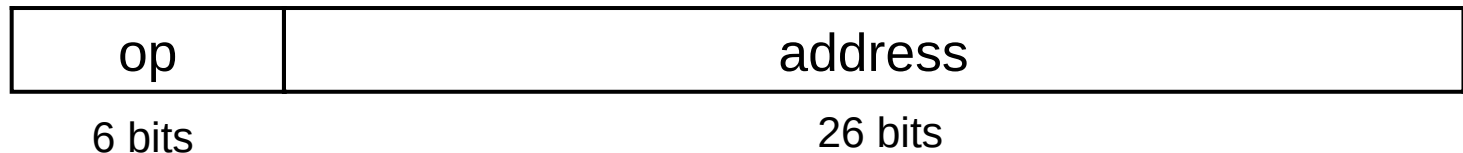
- (Pseudo)Direct jump addressing
    - Target address = $PC_{31\ldots28}$: (address × 4)

Target Addressing Example

- Loop code from earlier example
    - Assume Loop at location 80000

Loop: sll $t1, $s3, 2	80000	0	0	19	9	4	0

Loop: sll $t1, $s3, 2   80000

add $t1, $t1, $s6   80004

lw $t0, 0($t1)   80008

bne $t0, $s5, Exit   80012

addi $s3, $s3, 1   80016

j   Loop   80020

Exit: …   80024

80000	0	0	19	9	4	0
80004	0	9	22	9	0	32
80008	35	9	8	0		
80012	5	8	21	2		
80016	8	19	19	1		
80020	2		20000			
80024						

Branching Far Away

- If branch target is too far to encode with 16-bit offset, assembler rewrites the code

- Example

```
    beq $s0,$s1, L1
            ↓
    bne $s0,$s1, L2
    j L1
L2: …
```
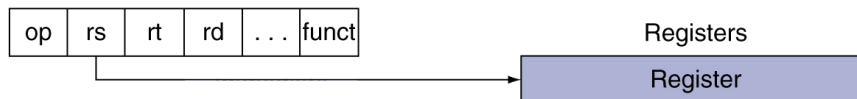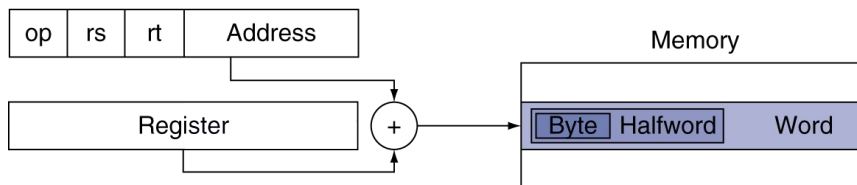
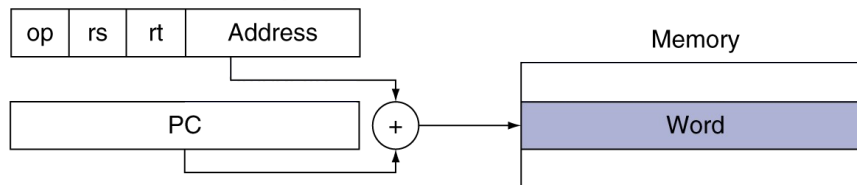# Addressing Mode Summary

1. Immediate addressing

| op | rs | rt | Immediate |

2. Register addressing

| op | rs | rt | rd | . . . | funct |

Registers

Register

3. Base addressing

| op | rs | rt | Address |

Register

+

Memory

Byte | Halfword | Word

4. PC-relative addressing

| op | rs | rt | Address |

PC

+

Memory

Word

5. Pseudodirect addressing

| op | Address |

PC

:

Memory

Word
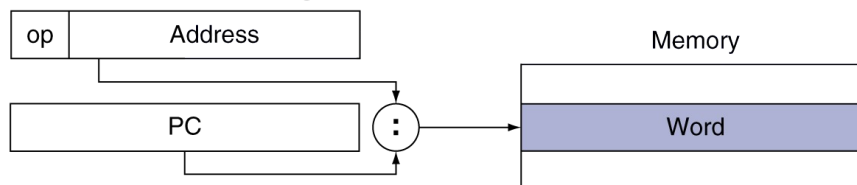
Register Usage

- $a0 – $a3: arguments (reg's 4 – 7)
- $v0, $v1: result values (reg's 2 and 3)
- $t0 – $t9: temporaries
  - Can be overwritten by callee
- $s0 – $s7: saved
  - Must be saved/restored by callee
- $gp: global pointer for static data (reg 28)
- $sp: stack pointer (reg 29)
- $fp: frame pointer (reg 30)
- $ra: return address (reg 31)