

# CS/ECE 3810: Computer Organization

## Lecture 4: MIPS instruction set

Anton Burtsev  
August, 2022

# Instruction Set

---

- Understanding the language of the hardware is key to understanding the hardware/software interface
- A program (in say, C) is compiled into an executable that is composed of machine instructions – this executable must also run on future machines – for example, each Intel processor reads in the same x86 instructions, but each processor handles instructions differently
- Java programs are converted into portable bytecode that is converted into machine instructions during execution (just-in-time compilation)
- What are important design principles when defining the instruction set architecture (ISA)?

# Instruction Set

---

- Important design principles when defining the instruction set architecture (ISA):
  - keep the hardware simple – the chip must only implement basic primitives and run fast
  - keep the instructions regular – simplifies the decoding/scheduling of instructions

We will later discuss RISC vs CISC

# A Basic MIPS Instruction

---

C code:  $a = b + c ;$

Assembly code: (human-friendly machine instructions)  
`add a, b, c    # a is the sum of b and c`

Machine code: (hardware-friendly machine instructions)  
`00000010001100100100000000100000`

Translate the following C code into assembly code:  
 $a = b + c + d + e ;$

# Example

---

C code `a = b + c + d + e;`  
translates into the following assembly code:

<code>add a, b, c</code>		<code>add a, b, c</code>
<code>add a, a, d</code>	or	<code>add f, d, e</code>
<code>add a, a, e</code>		<code>add a, a, f</code>

- Instructions are simple: fixed number of operands (unlike C)
- A single line of C code is converted into multiple lines of assembly code
- Some sequences are better than others... the second sequence needs one more (temporary) variable `f`

# Subtract Example

---

C code `f = (g + h) - (i + j);`

Assembly code translation with only add and sub instructions:

# Subtract Example

---

C code `f = (g + h) - (i + j);`  
translates into the following assembly code:

<code>add t0, g, h</code>		<code>add f, g, h</code>
<code>add t1, i, j</code>	or	<code>sub f, f, i</code>
<code>sub f, t0, t1</code>		<code>sub f, f, j</code>

- Each version may produce a different result because floating-point operations are not necessarily associative and commutative... more on this later

# Operands

---

- In C, each “variable” is a location in memory
- In hardware, each memory access is expensive – if variable *a* is accessed repeatedly, it helps to bring the variable into an on-chip scratchpad and operate on the scratchpad (registers)
- To simplify the instructions, we require that each instruction (add, sub) only operate on registers
- Note: the number of operands (variables) in a C program is very large; the number of operands in assembly is fixed... there can be only so many scratchpad registers



# Registers

---

- The MIPS ISA has 32 registers (x86 has 8 registers) – Why not more? Why not less?
- Each register is 32 bits wide (modern 64-bit architectures have 64-bit wide registers)
- A 32-bit entity (4 bytes) is referred to as a word
- To make the code more readable, registers are partitioned as \$s0-\$s7 (C/Java variables), \$t0-\$t9 (temporary variables)...

# Binary Stuff

---

- 8 bits = 1 Byte, also written as 8b = 1B
- 1 word = 32 bits = 4B
- 1KB = 1024 B =  $2^{10}$  B
- 1MB = 1024 x 1024 B =  $2^{20}$  B
- 1GB = 1024 x 1024 x 1024 B =  $2^{30}$  B
- A 32-bit memory address refers to a number between 0 and  $2^{32} - 1$ , i.e., it identifies a byte in a 4GB memory


# Memory Operands

---

- Values must be fetched from memory before (add and sub) instructions can operate on them


Load word

`lw $t0, memory-address`



Store word

`sw $t0, memory-address`




How is memory-address determined?

# Memory Address

---

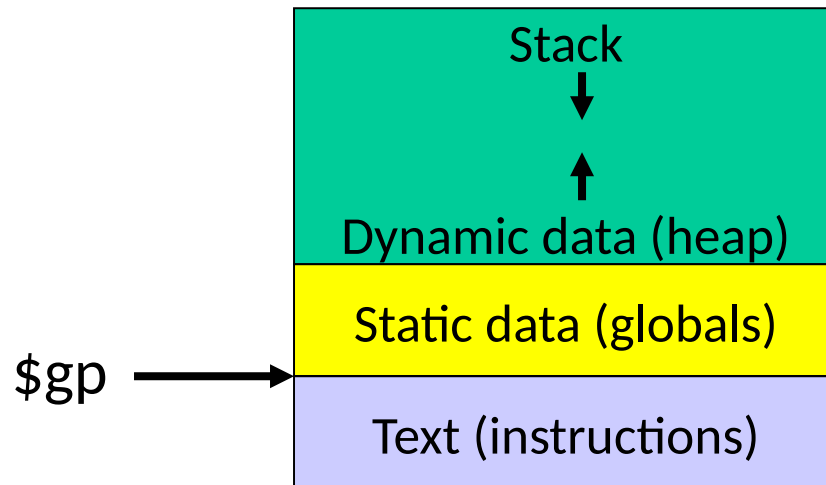
- The compiler organizes data in memory... it knows the location of every variable (saved in a table)... it can fill in the appropriate mem-address for load-store instructions



# Memory Organization

---

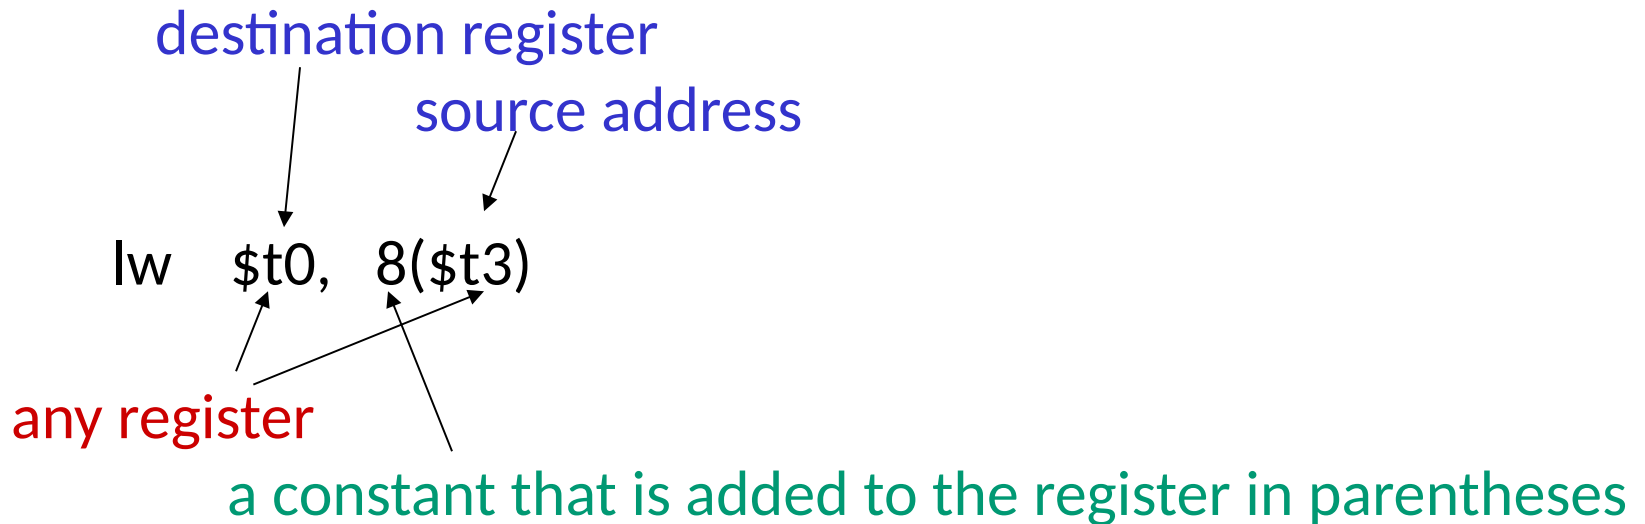
\$gp points to area in memory that saves global variables



# Memory Instruction Format

---

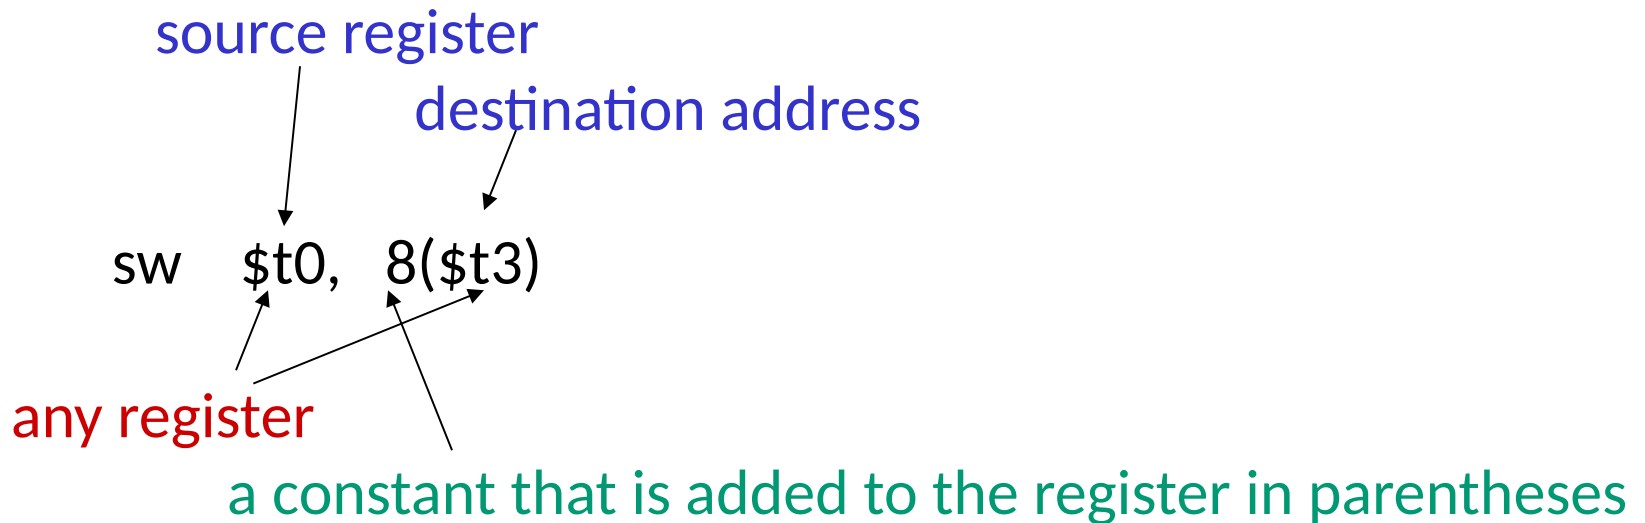
- The format of a load instruction:



# Memory Instruction Format

---

- The format of a store instruction:



# Example

---

```
int a, b, c, d[10];
```

```
addi $gp, $zero, 1000 # assume that data is stored at  
                        # base address 1000; placed in $gp;  
                        # $zero is a register that always  
                        # equals zero
```

```
lw  $s1, 0($gp) # brings value of a into register $s1  
lw  $s2, 4($gp) # brings value of b into register $s2  
lw  $s3, 8($gp) # brings value of c into register $s3  
lw  $s4, 12($gp) # brings value of d[0] into register $s4  
lw  $s5, 16($gp) # brings value of d[1] into register $s5
```



# Example

---

Convert to assembly:

C code: `d[3] = d[2] + a;`

# Example

---

Convert to assembly:

C code: `d[3] = d[2] + a;`

Assembly (same assumptions as previous example):

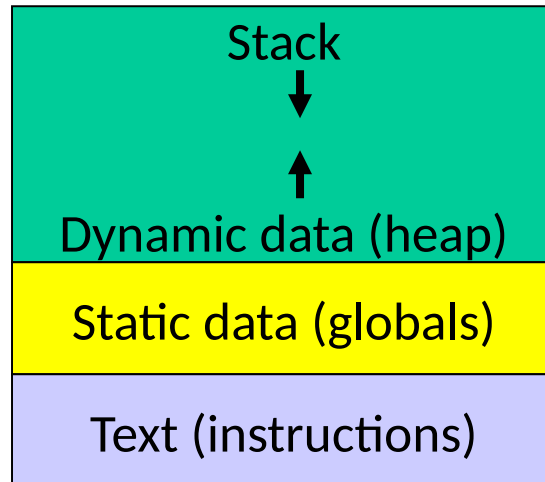
```
lw    $s0, 0($gp)    # a is brought into $s0
lw    $s1, 20($gp)   # d[2] is brought into $s1
add   $s2, $s0, $s1  # the sum is in $s2
sw    $s2, 24($gp)   # $s2 is stored into d[3]
```

Assembly version of the code continues to expand!

# Memory Organization

---

- The space allocated on stack by a procedure is termed the activation record (includes saved values and data local to the procedure) – frame pointer points to the start of the record and stack pointer points to the end – variable addresses are specified relative to \$fp as \$sp may change during the execution of the procedure
- \$gp points to area in memory that saves global variables
- Dynamically allocated storage (with malloc()) is placed on the heap



# Binary Representation

---

- The binary number

01011000 00010101 00101110 11100111

Most significant bit ←      ← Least significant bit

represents the quantity

$$0 \times 2^{31} + 1 \times 2^{30} + 0 \times 2^{29} + \dots + 1 \times 2^0$$

- A 32-bit word can represent  $2^{32}$  numbers between 0 and  $2^{32}-1$   
... this is known as the unsigned representation as we're assuming that numbers are always positive

# Negative Numbers

---

32 bits can only represent  $2^{32}$  numbers – if we wish to also represent negative numbers, we can represent  $2^{31}$  positive numbers (incl zero) and  $2^{31}$  negative numbers

$$0000\ 0000\ 0000\ 0000\ 0000\ 0000\ 0000\ 0000_{\text{two}} = 0_{\text{ten}}$$

$$0000\ 0000\ 0000\ 0000\ 0000\ 0000\ 0000\ 0001_{\text{two}} = 1_{\text{ten}}$$

...

$$0111\ 1111\ 1111\ 1111\ 1111\ 1111\ 1111\ 1111_{\text{two}} = 2^{31}-1$$

$$1000\ 0000\ 0000\ 0000\ 0000\ 0000\ 0000\ 0000_{\text{two}} = -2^{31}$$

$$1000\ 0000\ 0000\ 0000\ 0000\ 0000\ 0000\ 0001_{\text{two}} = -(2^{31} - 1)$$

$$1000\ 0000\ 0000\ 0000\ 0000\ 0000\ 0000\ 0010_{\text{two}} = -(2^{31} - 2)$$

...

$$1111\ 1111\ 1111\ 1111\ 1111\ 1111\ 1111\ 1110_{\text{two}} = -2$$

$$1111\ 1111\ 1111\ 1111\ 1111\ 1111\ 1111\ 1111_{\text{two}} = -1$$

# 2's Complement

0000 0000 0000 0000 0000 0000 0000 0000<sub>two</sub> = 0<sub>ten</sub>

0000 0000 0000 0000 0000 0000 0000 0001<sub>two</sub> = 1<sub>ten</sub>

...

0111 1111 1111 1111 1111 1111 1111 1111<sub>two</sub> =  $2^{31}-1$

1000 0000 0000 0000 0000 0000 0000 0000<sub>two</sub> =  $-2^{31}$

1000 0000 0000 0000 0000 0000 0000 0001<sub>two</sub> =  $-(2^{31} - 1)$

1000 0000 0000 0000 0000 0000 0000 0010<sub>two</sub> =  $-(2^{31} - 2)$

...

1111 1111 1111 1111 1111 1111 1111 1110<sub>two</sub> = -2

1111 1111 1111 1111 1111 1111 1111 1111<sub>two</sub> = -1

Why is this representation favorable?

Consider the sum of 1 and -2 .... we get -1

Consider the sum of 2 and -1 .... we get +1

This format can directly undergo addition without any conversions!

Each number represents the quantity

$$x_{31} - 2^{31} + x_{30} 2^{30} + x_{29} 2^{29} + \dots + x_1 2^1 + x_0 2^0$$

# 2's Complement

0000 0000 0000 0000 0000 0000 0000 0000<sub>two</sub> = 0<sub>ten</sub>  
0000 0000 0000 0000 0000 0000 0000 0001<sub>two</sub> = 1<sub>ten</sub>  
...  
0111 1111 1111 1111 1111 1111 1111 1111<sub>two</sub> =  $2^{31}-1$   
  
1000 0000 0000 0000 0000 0000 0000 0000<sub>two</sub> =  $-2^{31}$   
1000 0000 0000 0000 0000 0000 0000 0001<sub>two</sub> =  $-(2^{31} - 1)$   
1000 0000 0000 0000 0000 0000 0000 0010<sub>two</sub> =  $-(2^{31} - 2)$   
...  
1111 1111 1111 1111 1111 1111 1111 1110<sub>two</sub> = -2  
1111 1111 1111 1111 1111 1111 1111 1111<sub>two</sub> = -1

Note that the sum of a number  $x$  and its inverted representation  $x'$  always equals a string of 1s (-1).

$$x + x' = -1$$

$x' + 1 = -x$  ... hence, can compute the negative of a number by

$-x = x' + 1$  inverting all bits and adding 1

Similarly, the sum of  $x$  and  $-x$  gives us all zeroes, with a carry of 1

In reality,  $x + (-x) = 2^n$  ... hence the name 2's complement

# Example

---

- Compute the 32-bit 2's complement representations for the following decimal numbers:  
5, -5, -6



# Example

---

- Compute the 32-bit 2's complement representations for the following decimal numbers:  
5, -5, -6

5: 0000 0000 0000 0000 0000 0000 0000 0101

-5: 1111 1111 1111 1111 1111 1111 1111 1011

-6: 1111 1111 1111 1111 1111 1111 1111 1010

Given -5, verify that negating and adding 1 yields the number 5

# Signed / Unsigned

---

- The hardware recognizes two formats:

unsigned (corresponding to the C declaration `unsigned int`)

-- all numbers are positive, a 1 in the most significant bit just means it is a really large number

signed (C declaration is `signed int` or just `int`)

-- numbers can be +/- , a 1 in the MSB means the number is negative

This distinction enables us to represent twice as many numbers when we're sure that we don't need negatives

# MIPS Instructions

---

Consider a comparison instruction:

`slt $t0, $t1, $zero`

and \$t1 contains the 32-bit number `1111 01...01`

What gets stored in \$t0?

# MIPS Instructions

---

Consider a comparison instruction:

`slt $t0, $t1, $zero`

and `$t1` contains the 32-bit number `1111 01...01`

What gets stored in `$t0`?

The result depends on whether `$t1` is a signed or unsigned number – the compiler/programmer must track this and accordingly use either `slt` or `sltu`

`slt $t0, $t1, $zero` stores 1 in `$t0`

`sltu $t0, $t1, $zero` stores 0 in `$t0`

# Sign Extension

---

- Occasionally, 16-bit signed numbers must be converted into 32-bit signed numbers – for example, when doing an add with an immediate operand
- The conversion is simple: take the most significant bit and use it to fill up the additional bits on the left – known as sign extension

So  $2_{10}$  goes from 0000 0000 0000 0010 to  
0000 0000 0000 0000 0000 0000 0000 0010

and  $-2_{10}$  goes from 1111 1111 1111 1110 to  
1111 1111 1111 1111 1111 1111 1111 1110

# Alternative Representations

---

- The following two (intuitive) representations were discarded because they required additional conversion steps before arithmetic could be performed on the numbers
  - sign-and-magnitude: the most significant bit represents +/- and the remaining bits express the magnitude
  - one's complement:  $-x$  is represented by inverting all the bits of  $x$

Both representations above suffer from two zeroes